

Master 08 — Developer RFP / Intervention Test Build Specification

Version: v1.3 — updated 22 July 2026, 22:40 BST. Versioning scheme confirmed by `ProjectFiles.md §14` (major.minor; minor bump = clarification/correction added, central rule unchanged) — see change log at the end of this document.

Audience: Technical suppliers, product architects, security / IG reviewers, OTOS technical advisors.

Purpose: Specify the demo build and future-platform requirements while preventing clinical-record, EPR or automated-triage drift.

Status: PENDING CLINICAL / GOVERNANCE REVIEW.

Priority: P1.

Review needed from: Dean; technical lead; IG / DPO; security reviewer; governance lead; supplier.

Source basis: Master 08 source text, Programme Plan technical build model, data model, START replacement evidence / data section.

A4 REVIEW PAGE 1 — Purpose and non-negotiables

This RFP seeks serious technical proposals for OTOS Continuity™ as a health-adjacent, non-clinical continuity infrastructure platform.

The document must keep two scopes separate:

1. **Intervention Test build** — the first bounded build for up to 50 adults. This is the immediate supplier ask.
2. **Full Pilot / future platform** — a separate later scope, priced and governed separately.

The first build must support a bounded intervention test. It must not accidentally become a clinical system, a pathway-management system or a full pilot platform.

Non-negotiables:

- no diagnosis;
- no prescribing;
- no medication management;
- no treatment decisions;
- no EPR write-back;
- no clinical-note storage;
- no therapy-note storage;
- no automated clinical triage;
- human-in-control review;
- tokenised identity;
- consent capture;
- audit-grade receipts;
- UK-hosted infrastructure;
- role-based access;
- SMS Engine designed properly from the start **CLARIFIED BELOW** ;
- evidence export for the formal decision set.

New in v1.1 — SMS Engine vs SMSC, scope clarification. "SMS Engine designed properly from the start" means an **application layer** that integrates with a third-party SMS gateway/aggregator provider — it does not mean building carrier-grade telecom infrastructure (an SMSC). An SMSC is the store-and-forward system mobile network operators themselves run, with direct SS7/SIGTRAN interconnects into every UK mobile network; building one is equivalent to becoming a telecoms carrier — Ofcom-level obligations, direct commercial agreements with mobile networks, carrier-grade redundancy. That is wildly disproportionate to a bounded test for up to 50 adults and is **not** part of this ask.

What the supplier should price and build: an SMS Engine that sends partner-approved templated messages via a chosen aggregator's API (e.g. Twilio, MessageBird, Sinch, Esendex — all existing, UK-hosted, already carrier-connected); receives inbound replies/opt-outs via webhook; captures delivery receipts; handles retries, quiet-hours and rate limits; and logs every event as `MessageEvent` / `SmsDeliveryReceipt` on the Thread.

Aggregator capacity is effectively unlimited for this test's purposes — the interaction-frequency question (how many touches per participant per day/week) is therefore a **design and cost decision**, not a technical ceiling, and suppliers should be asked to quote against an assumed cadence rather than asked to size for a capacity limit. Message content, the non-clinical wording rule, and the full audit trail remain fully under OTOS's own control regardless of which aggregator is chosen — the aggregator only moves the message, it never decides content or holds the record.

OTOS exists to test whether adults whose addiction, relapse or repeated crisis may be driven by untreated or inadequately treated ADHD can be identified, clinically routed, supported through authorised partner continuity activity, followed through a whole-person Continuity Thread, and evidenced safely.

The first build must create the **Continuity Management Layer** and the user's **Continuity Thread** — not a pathway. The Layer is the tokenised, role-based, consent-led infrastructure partners and the user access. The Thread is the whole-person journey over time, built from agreed non-clinical events, touchpoints, handoffs, prompts, responses, drift / silence signals, review actions and evidence-export data. Partners keep their own service routes, records, responsibilities and clinical communication. OTOS provides the infrastructure for authorised partners to view, message, log, alert and manage their agreed part of continuity without merging service records or shifting clinical responsibility.

New in v1.2 — partner connection model and anti-pass-through rule. Each partner's own existing service route to the user (their own records, communications, responsibilities) stays completely separate from OTOS and continues directly between that service and the user, exactly as it does today. The Continuity Management Layer sits alongside those routes, not inside them — it is one shared, externally managed, non-clinical layer, accessed separately by each partner through their own login.

The sensitive point the build must get right: a partner submits agreed touchpoint/interaction data *into* the Layer, and can only see back what's agreed for their own part of that user's Continuity Thread — never another partner's data, and never another partner's interface. **Partners do not connect to one another through OTOS, at any point.** Internally, the developer may find it natural to implement each partner's connection as its own distinct object on the network (a "Partner Dedicated Node" or equivalent) so this data boundary is enforced structurally, not just by convention — but that's an implementation detail for the developer to name and build correctly, not a fixed spec term OTOS is mandating. What's non-negotiable is the boundary itself: one Layer, one Thread per user, each partner's own separate access, no partner-to-partner visibility, existing service routes untouched and running in parallel.

A4 REVIEW PAGE 2 — Required response format

Suppliers should structure responses using these headings:

1. Executive technical summary.
2. Assumptions and interpretation of scope.
3. Proposed architecture.
4. Intervention Test build scope and cost.
5. Full Pilot / future platform scope and cost.
6. SMS Engine / communications architecture — see the SMS Engine vs SMSC clarification on page 1; suppliers should describe their chosen aggregator integration, not carrier-level infrastructure.
7. Security, encryption and audit model.
8. Data model.
9. Role-based access and tenancy model.
10. Delivery plan and milestones.
11. Person-days by role.
12. Commercials: non-recurring engineering, run costs and support.
13. Risks, dependencies and exclusions.
14. Similar work / team credentials.
15. Questions for OTOS before final quote.

Suppliers must explicitly state where they believe a requirement creates clinical, IG, security, safeguarding, cost or delivery risk.

A4 REVIEW PAGE 3 — Scope A and Scope B

Scope A — Intervention Test build

The Intervention Test build tests whether a consent-led Continuity Management Layer can be provided around a small, governed user cohort, with a Continuity Thread following each user's journey over time, and whether this can be evidenced safely without creating a new clinical pathway or shared clinical record.

It should support up to 50 adults, with sensible buffer capacity if agreed.

Primary purpose:

- test whether continuity receipts work;
- test whether micro mobile check-ins can keep contact warm;
- test whether SMS delivery / failure / reply / silence can be captured;
- test whether drift / silence can be surfaced earlier;
- test whether human review can be auditable;
- test whether an evidence export can be produced;
- test whether the Continuity Management Layer and user Thread can be governed without EPR integration;
- test whether partner interfaces / authorised views can access the Layer to log touchpoints and view agreed Thread elements without merging service records;
- test whether partner-approved messages, prompts and check-ins can be delivered through the Continuity Management Layer and recorded as agreed events on the user's Continuity Thread.

This is not the full platform and not the Full Pilot.

Scope B — Full Pilot / future platform specification

The Full Pilot / future platform specification is separate from the Intervention Test build. It may include: multiple partners; configurable user journeys / service-route patterns; multiple touchpoint types; partner-specific dashboards; communications-provider abstraction; consent and preference management; real-world evidence outputs; more sophisticated analytics; operational governance; future integration readiness.

Scope B must be priced separately. It is not assumed inside the Intervention Test build budget.

A4 REVIEW PAGE 4 — Demo build requirements

The Intervention Test build should support: participant tokenisation; consent capture; withdrawal mechanism; minimal personal data; contact method for SMS; communication preferences; partner / source service-line field; ADHD clinical-route status / ADHD next-step status field where agreed; clinical-owner field; medication-status category where agreed; addiction-service touchpoints; structured support / waiting-well touchpoints; hospital checkpoint where agreed; handoff receipts; partner-approved SMS check-in schedules; partner-approved outbound SMS; inbound SMS where possible; delivery receipts; opt-out; message event log; inbound response handling; drift / silence rules; human-review queue; review decision log; dashboard; evidence export; audit log; role-based access.

The participant-token link should remain with the referring service wherever possible. OTOS should not require NHS number, date of birth or full clinical identity for the first test unless governance explicitly requires it.

New in v1.3 — partner-by-partner First-50 interaction scope. Not every partner has the same footprint in the Intervention Test build. This is a genuine build-scope decision, not a diagram styling choice — the data model and dashboard must reflect these differences, not treat all four partners as equally rich from day one.

Partner	First-50 scope	Why
RCE / Waiting-Well	Full functionality, live from the start. Frequent partner-approved micro-check-ins and course-tied SMS prompts, replies, attendance, silence/drift signals — the richest of the four.	This isn't a future add-on. It's core to how the Intervention Test actually produces its results and evidence, and to the case for moving to a Full Pilot.
CGL / CRS	Moderate-to-rich. Multiple touchpoint/message types beyond the basics (e.g. attendance confirmation plus partner-specific prompts such as a relapse-prevention check-in) — exact set to be agreed per partner.	CGL/CRS keyworkers have said this kind of reinforcement is useful to them regardless of OTOS — see the Science of Sticking rationale.
ADHD Team	Minimal, possibly a single touchpoint type (e.g. "attended: assessment") — not yet confirmed.	Genuinely unknown until confirmed directly with the team. Build and diagrams should show this partner deliberately light, not padded out to match the others.
Alcohol Liaison	Acknowledgement-only in the First-50. A single flag if one of the 50 presents through A&E and onto a ward for detox. No onboarding, no signposting, in this test.	Deliberately narrow scope for Phase 1. Full Pilot direction: this becomes full bedside onboarding of new users at the point of crisis admission — potentially one of the biggest upstream intervention moments in the whole model. First-50 tests the infrastructure; it doesn't yet test that bedside-onboarding moment.

Any diagram or partner close-up showing these four must reflect this variance accurately — RCE richest, Alcohol Liaison narrowest with a clear forward-note toward its Full Pilot bedside-onboarding role, ADHD Team kept light pending confirmation, CGL/CRS in between.

A4 REVIEW PAGE 5 — Evidence model: thread signal, not clinical record

The build must evidence agreed non-clinical events in the user's Continuity Thread and service journey. It must not create a clinical record, shared EPR or new pathway.

START strengthens the need to monitor the ADHD-addiction sequence together, but the system must remain non-clinical.

The build must capture: whether agreed thread events happened; whether communication landed; whether silence or drift emerged; whether human review occurred; whether support and addiction touchpoints were visible; whether a Week-12 evidence export can support a continue / continue with changes / rebuild / pause / stop decision.

The build must not store: clinical notes; therapy notes; prescribing instructions; diagnostic records; medication dose advice; full service records; EPR content; private session content; automated clinical-risk decisions.

The Continuity Management Layer is shared continuity infrastructure, not a shared clinical record. The Continuity Thread follows the user journey through agreed non-clinical events.

A4 REVIEW PAGE 6 — Minimum data model

The build should support these minimum entities or equivalent structures:

Entity	Purpose
ParticipantToken	Tokenised continuity subject
ConsentRecord	Consent and withdrawal
ContactMethod	SMS contact route
CommunicationPreference	Cadence, opt-out and template permissions
Cohort	Demonstrator cohort
PartnerOrganisation	Partner / service line
StaffUser	Staff access
Touchpoint	Service or support contact
HandoffReceipt	Next-step receipt
CheckInSchedule	Partner-approved planned prompts
CheckInResponse	Participant reply
DriftSignal	Silence / stuck signal
HumanReview	Review workflow
ReviewDecision	Decision log
MessageTemplate	Partner-approved SMS text
MessageEvent	SMS audit
SmsDeliveryReceipt	Provider receipt
EvidenceExport	Reporting pack
AuditEvent	Append-only audit

The supplier may propose alternative entity names, but must preserve the boundary: continuity infrastructure and Thread events only, no clinical record.

A4 REVIEW PAGE 7 — Dashboard and evidence export

The dashboard is a role-based view of agreed Layer activity and authorised Thread events. It is not a clinical dashboard and must not imply full visibility across partner records.

The first dashboard should include: participants onboarded; active participants; consent complete / withdrawn; handoff receipts created; handoffs confirmed as landed; open drift signals; human reviews completed; median time to review; SMS sent / delivered / failed / replied / opted out; check-in completion rate; unresolved items; partner burden indicator.

Required evidence export categories: cohort summary; consent / withdrawal summary; ADHD clinical-route / next-step status summary where agreed; medication-status category summary where agreed; addiction-service touchpoint report; structured support / waiting-well touchpoint report; handoff receipt report; SMS delivery / failure / reply / silence report; drift / silence report; human-review report; partner burden report; safety / incident report; evidence-labelled claims table; continue / continue with changes / rebuild / pause / stop recommendation.

A4 REVIEW PAGE 8 — Security, IG and exclusions

The supplier must describe: hosting location; encryption at rest; encryption in transit; secrets management; role-based access; audit trail; user provisioning and de-provisioning; backup and restore; incident response; data retention and deletion support; SMS provider security; sub-processor position; support model; monitoring; logging; vulnerability management.

The first build should not require: EPR integration; NHS number matching; clinical-note storage; prescribing workflow; diagnostic workflow; AI clinical triage; full multi-tenant platform; national rollout architecture; **any partner-to-partner data access or connection** — each partner's access must be structurally isolated to their own agreed part of the Continuity Thread.

The first build should provide the Continuity Management Layer for the intervention test only; it must not become a full pathway-management platform by accident. It should be built in a way that does not block future scale if the evidence supports continuing.

Review notes

Dean decision needed

~~[RESOLVED] Confirm whether the RFP asks for demo build only or demo plus future platform pricing.~~

Dean note: Ask suppliers to separate demo build pricing from future platform pricing. Scope A and Scope B must stay separate.

~~[RESOLVED] Confirm whether the technical brief separates the Intervention Test from the Full Pilot.~~

Dean note: Yes. The immediate supplier ask must be the Intervention Test build. The Full Pilot / future platform is a separate later scope, priced and governed separately. Do not let the first build become a full pilot, pathway-management system or clinical platform.

~~[RESOLVED] Confirm whether buffer capacity beyond up to 50 adults remains in the technical brief.~~

Dean note: Keep sensible buffer capacity if it does not distort the first-test scope, but the first demonstrator remains up to 50 adults.

~~[RESOLVED] Confirm whether SMS provider preference is known or left open.~~

Dean note: Leave SMS provider open for supplier recommendation unless governance, cost, UK hosting, deliverability or procurement requirements dictate a preferred provider. (22 Jul update: confirmed open — provider is an aggregator choice, not an OTOS-built SMSC; see page 1.)

~~[RESOLVED] Confirm whether supplier should price evaluation support or export only.~~

Dean note: Supplier should price evidence export clearly. Evaluation support can be optional / separately priced so it does not blur technical delivery with evaluation ownership.

[OPEN] Confirm assumed interaction cadence (touches per participant per day/week) for suppliers to quote against.

New open item, 22 Jul 2026: Not yet set. This is a design/cost decision, not a technical ceiling (see page 1 clarification) — needs a working number before the RFP goes to suppliers so they can size SMS volume and cost.

Evidence / citation check needed

[RESOLVED] No evidence citation is needed for the technical procurement draft unless evidence claims are retained in supplier-facing copy.

Dean note: Tick. No evidence citation is needed for the technical procurement draft unless evidence claims are retained in supplier-facing copy.

Clinical / governance check needed

[OPEN] IG / DPO review required before supplier circulation.

Dean note: Leave unticked for now. This needs IG / DPO input before approval.

[OPEN] Security review required.

Dean note: Leave unticked for now. This needs security / technical-governance input before approval.

[OPEN] Clinical boundary review required.

Dean note: Leave unticked for now. This needs clinical / governance input before approval.

[OPEN] Confirm medication-status category fields.

Dean note: Leave unticked for now. This needs clinical-team input before approval.

[OPEN] Confirm human-review workflow and safeguarding escalation fields.

Dean note: Leave unticked for now. This needs safeguarding / governance input before approval.

Copy approval

☐ Approved ☐ Needs edit ☐ Hold

Change log

Date / time	Change	By
9 Jul 2026 (original)	Master 08 created — Developer RFP / Intervention Test Build Specification, v1.	OTOS build process (pre-Cowork)
22 Jul 2026, 19:46 BST	Added SMS Engine vs SMSC scope clarification (page 1 and response-format section 6); added open review item for interaction-cadence decision; no other content changed. Superseded file: Master 08 - Developer RFP Intervention Test Build.pdf (v1, preserved unchanged for provenance).	Claude, on Dean's direct instruction
22 Jul 2026, 21:05 BST	Added the partner connection model / anti-pass-through rule (page 1 callout and a new exclusion on page 8), to bring this document in line with the Layer/Thread/Partner-Node architecture locked earlier the same day in the diagram-correction work and restated in OTOS_HUB_TRUTH_PROTOCOL.md §8. This closes the gap identified when Dean proposed OTOS_ARCHITECTURE_AUTHORITY.md, which names this RFP as the absolute build authority — that only holds if the RFP itself carries the current locked architecture, which it didn't until this edit. "Partner Dedicated Node" is presented here as an illustrative/ implementation concept for the developer, not a mandated spec term, per Dean's direct clarification. Superseded file: Master 08 - Developer RFP Intervention Test Build v1.1.pdf (preserved unchanged for provenance).	Claude, on Dean's direct instruction
22 Jul 2026, 22:40 BST	Added the partner-by-partner First-50 interaction scope table (new page 4a): RCE/Waiting-Well full functionality from the start; CGL/CRS moderate-to-rich; ADHD Team minimal/unconfirmed; Alcohol Liaison acknowledgement-only (A&E-to-detox-ward flag), explicitly not onboarding or signposting in the First-50, with a forward-note to its Full Pilot bedside-onboarding role. Raised directly by Dean during the diagram-brief conversation — this is genuine build scope, not a visual styling choice, so it belongs in the RFP rather than only in a diagram brief. Superseded file: Master 08 - Developer RFP Intervention Test Build v1.2.pdf (preserved unchanged for provenance).	Claude, on Dean's direct instruction

This is a working draft (v1.3). Versioning follows ProjectFiles.md §14. Prior versions (v1, v1.1, v1.2) preserved unchanged alongside this version for provenance.

Appendix note — visual architecture reference

The technical drawing brief has been moved out of this supplier-facing RFP and should be held as a separate visual / diagram-production reference: see **Appendix — Technical Drawing Brief / Visual Architecture Reference v2** (updated 22 Jul 2026 alongside this document).

Use the Developer RFP above as the build brief. Use the separate visual architecture reference only to create or check diagrams that explain the build.